

University of Pennsylvania
Department of Electrical and System Engineering
System-on-a-Chip Architecture

ESE5320, Fall 2023 Design and Function Milestone Wednesday, November 1

Due: Friday, Nov. 10, 5:00PM

Group: Develop functional code. Identify design space options. Writeup (single turn-in for group).

1. Identify major design space axes that could be explored for your implementation.
 - For this milestone, aim for breadth (quantity of options)
 - Each axis description can be 1–3 sentences. Identify challenge being addressed, basic solution opportunity, and continuum. A single point in the design space is not a continuum; except in rare cases, this should capture a range of potential parameter values.
 - Include a simple equation to illustrate ideal benefit (e.g., running N tasks in parallel reduces runtime by a factor of N ; $T(N) = T(1)/N$) and the associated resource costs.
 - Cover all operations except SHA that must be accelerated including communication among operators. (i.e., CDC, Deduplication, LZW, and communication/integration)
 - You have four top-level options for SHA as noted on the next page which will make some of your primary design choices there, so we do not ask you to detail SHA here, but you should include communication with the SHA operations.
 - Aim for at least 6 axes per operation. Identify a few associated with how operations interact with each other.
 - Some of this should build on the parallelism opportunities you identified on the previous milestone.

Example from FFT design discussed in class.

Axis:	P , number of butterfly units.
Challenge:	Improving the throughput of the FFT
Opportunity:	Implement multiple hardware butterfly datapath units.
Continuum:	This can range from 1 to a fully spatial design with $P = \frac{N}{2} \log(N)$ butterfly units.
Equation for Benefit:	$Throughput(P) = P \times SingleButterflyThroughput$
Equation for Resources:	$Resources(P) = P \times SingleButterflyResources$

2. Refine your placeholder implementation into a functional implementation for the project task that can run on a Zynq ARM Cortex A53 processor and produce a valid compressed output stream that works with the supplied decompressor. Integrate with the provided ethernet input flow. Compress from ethernet input to SDCard output.
 - The primary goal for this assignment is functionality. As such, you should focus on a simple design that captures the necessary behavior.
 - As a result, this design need not be efficiently synthesizable to hardware.
 - However, you will eventually be optimizing this design and likely exploring HLS mappings to hardware. So, given a choice, you might want to use design constructs and idioms that you know will be more amenable to HLS hardware mappings.
 - Alternately, you should be prepared to rewrite your code later for efficient hardware mappings.
3. Turn in a tar or zip file with your functional code to the designated assignment component in canvas.
4. Turn in a tar or zip file with binaries to support execution of your code to the designated assignment component in canvas.
 - (a) The tar (or zip) files should include:
 - **encoder** – binary for your encoder to run on the the Ultra96
 - (b) Your encoder should take one argument:
 - the file name where the program should store the compressed data.
5. Measure the raw ethernet performance from your host machine to your Ultra96 (See Section 13.1 in Project Handout).
6. Document your design.
 - (a) Code sources (e.g., URLs) for any open-source code you used as a starting point or as a primary reference
 - (b) Current compression ratio and breakdown of contribution from deduplication and from LZW compression.
 - (c) Overall throughput (Gb/s) of your current implementation.
 - (d) Description of all validation performed on your current functional implementation.
 - (e) Report the raw ethernet speed measurements (Problem 5).
 - (f) Description of who did what. How did your team collaborate on the design, implementation, and validation?

7. Identify any challenges your group had in collaboration and design integration this week and how you plan to address them for future weeks.
 - This is not a question about technical status – that should be addressed above.
 - This is for teamwork, coordination, collaboration, communication, and workflow issues.
 - These may be things you’ve overcome by submission but didn’t go as smoothly as they should have.
 - In the unlikely case that everything went perfectly, identify the things you did that made it work well. For your future plans, look forward to next week to see if the same techniques are applicable or if there are new challenges that might require different or additional techniques for things to continue to go well.

SHA

You have four options for implementing SHA in your encoder pipeline:

1. Writing serialized SHA-256 or SHA3-384 running on the ARM processor.
2. Using the dedicated SHA3-384 unit in the Zynq Ultrascale.
3. Using SHA3-384 NEON intrinsics.
4. Implementing SHA-256 or SHA3-384 on FPGA using HLS (not until P3).

The specific option you choose will have implications on the maximum throughput you can achieve for your encoder pipeline. Moreover, you will need to adjust your design accordingly if you end up using the SHA3 unit, since it has a different digest size.

- You can find throughput comparisons of SHA3-384 on NEON and the dedicated SHA3 unit here: https://www.xilinx.com/support/documentation/white_papers/wp512-accel-crypt.pdf.
- You can find an example of how to use the SHA3 unit here: <https://xilinx-wiki.atlassian.net/wiki/spaces/A/pages/18841654/Linux+SHA+Driver+for+Zynq+Ultrascale+MPSoc> Note we have already enabled the SHA3 unit and user-space API in a provided platform. You will just need to write the driver.
 - *N.B.* To use this unit, data provided must be padded to a multiple of 4B to get consistent results; providing this padding and extracting full performance can be tricky.
- You can find software implementations that use NEON intrinsics here:
 - <https://github.com/james-ben/mpsoc-crypto>
 - <https://github.com/nolader/SHA-Intrinsics>

Multiple Cores

Recall from Homework 3 that you can utilize multiple cores using `std::threads`. There are four cores (ARM Cortex-A53) on the Ultra96. Moreover, recall that each core on the Ultra96 has one 64-bit NEON SIMD unit with 128-bit registers that you can utilize simultaneously with `std::threads`.

FPGA Acceleration Tutorial: Bloom Filter

Using bloom filter as the application, this tutorial shows you:

- a significant speedup ($8\times$) when computations are offloaded on the FPGA efficiently.
- how to write an HLS kernel for a CPU implementation (using `ap_uint`, `hls::stream`, and `pragmas`).
- how to achieve communication-compute overlap using sub-buffers.

1. Clone the `ese532_code` repository using the following command:

```
git clone https://github.com/icgrp/ese532_code.git
```

If you already have it cloned, pull in the latest changes using:

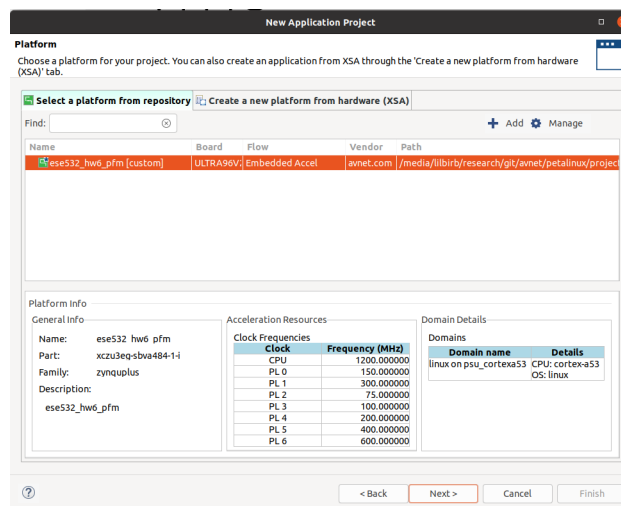

```
cd ese532_code/  
git pull origin master
```

The code you will use for this section is in the `vitis_tutorials/bloom` directory. The directory structure looks like this:


```
bloom/  
  cpu/  
  fpga/  
    MurmurHash2.c  
    common.h  
    compute_score_fpga_kernel.cpp  
    compute_score_host.cpp  
    hls_stream_utils.h  
    main.cpp  
    sizes.h  
    xcl2.cpp  
    xcl2.hpp
```

The `cpu` folder has a standalone CPU implementation of the bloom filter, which you can compile using the Vitis GUI flow from P2 and run it. The `main.cpp` code in the `fpga` folder has the OpenCL host code. The top level HLS function is in `compute_score_fpga_kernel.cpp`. We will now show how to use the Vitis GUI flow to compile OpenCL and HLS code.

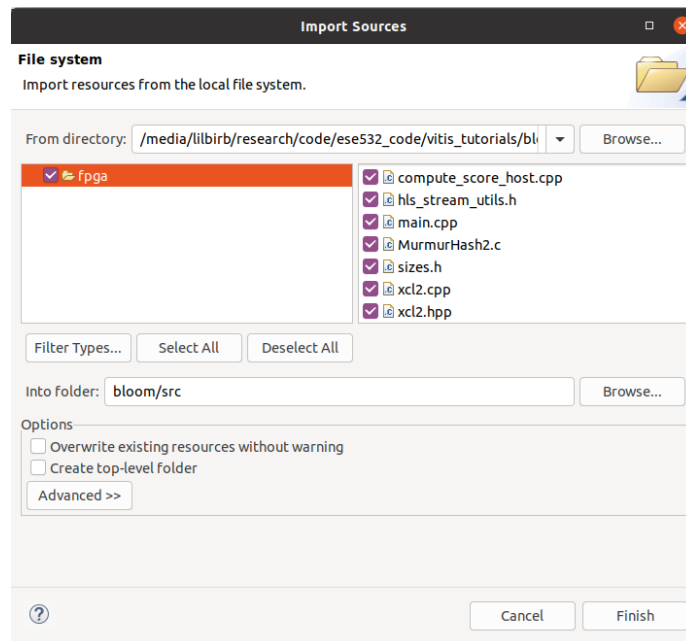
2. Create or use an existing workspace, create a new application project and use the provided platform.




3. Choose Empty Application from the SW acceleration templates as follows:



- Right-click on the `src` folder and click on import sources. Import the source files for this project as follows:

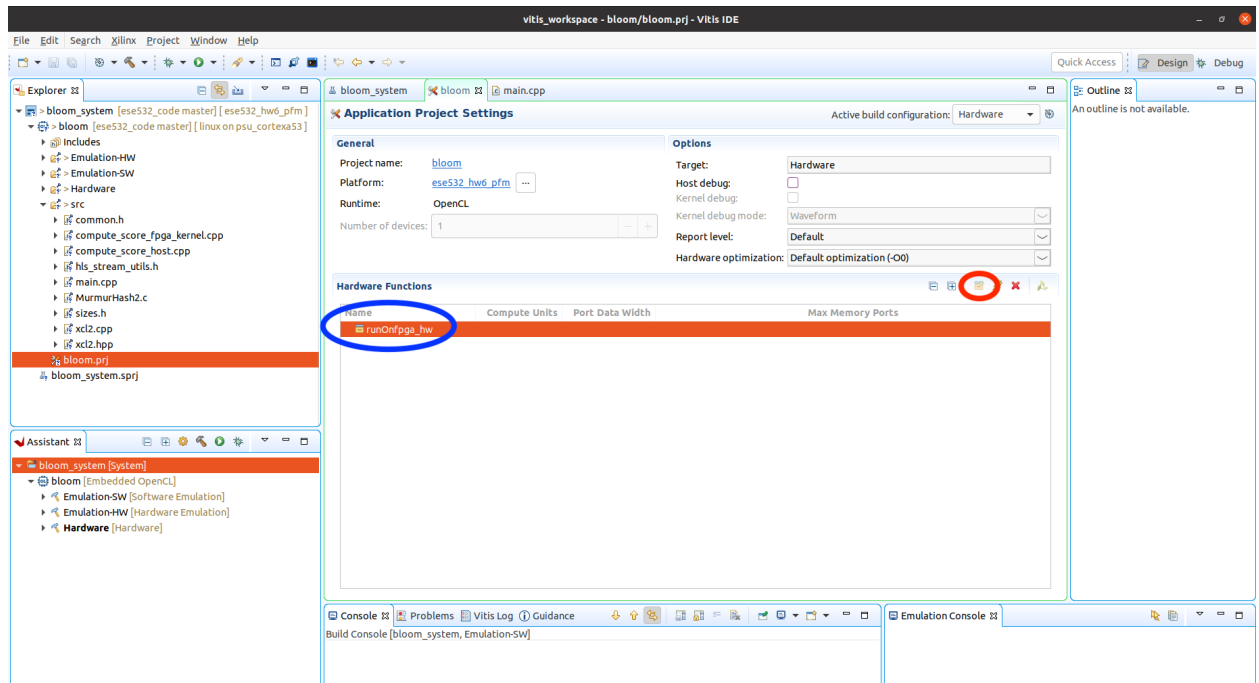


- Click on `bloom.prj` from the Explorer and change the Active build configuration to **Hardware** as shown below:




- Add an `xclbin` container by clicking on the button circled in red below. Rename the


container to `runOnfpga_hw` by clicking on the name `binary_container_1`. This is the `xclbin` name that the host code uses.



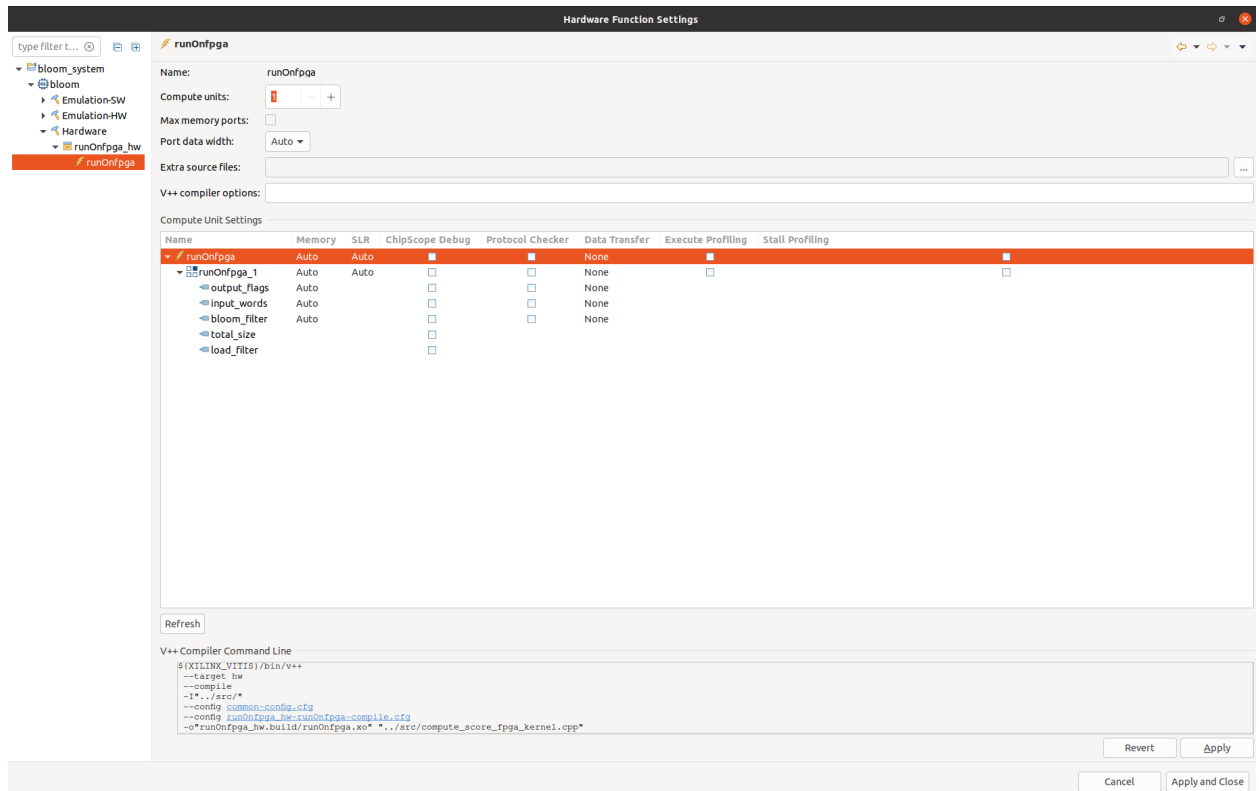
- Now click on the button circled in blue below. Choose the `runOnfpga` function as the hardware function.




8. From the Assistant view, double click on Hardware as follows:



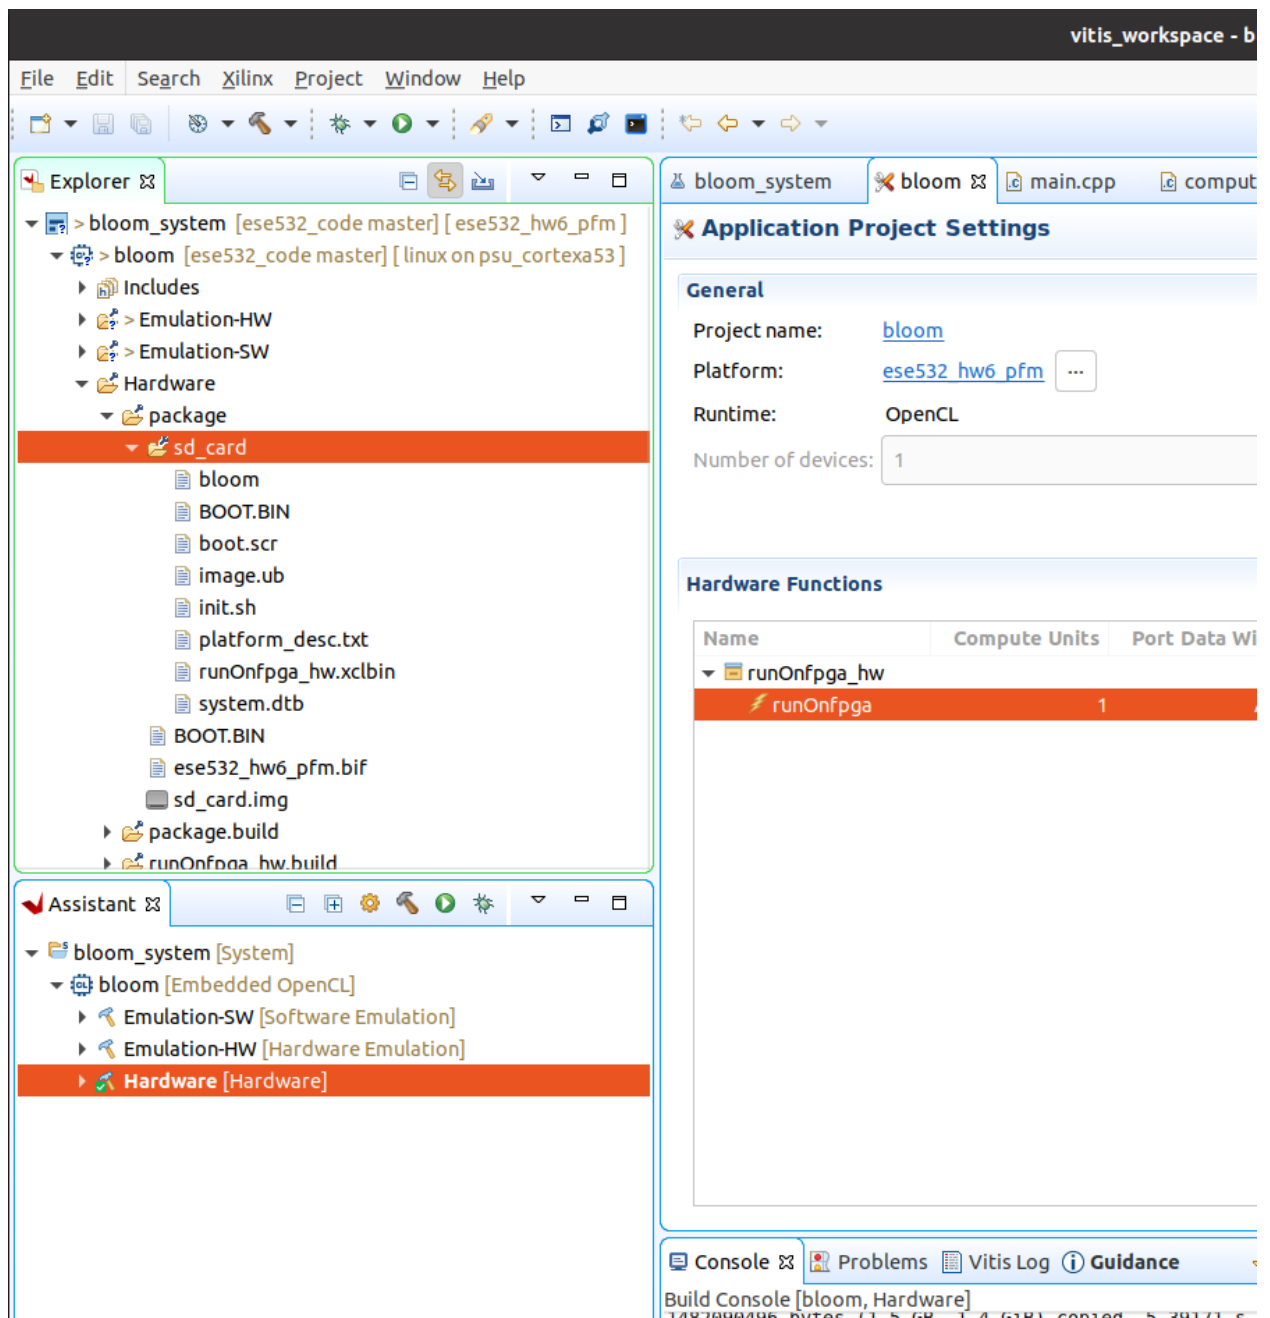
- Click on Hardware→runOnfpga.hw→runOnfpga. This brings the screen where you can specify the number of compute units, compiler options for the kernel, assign different ports to inputs etc. Keep the defaults for now:



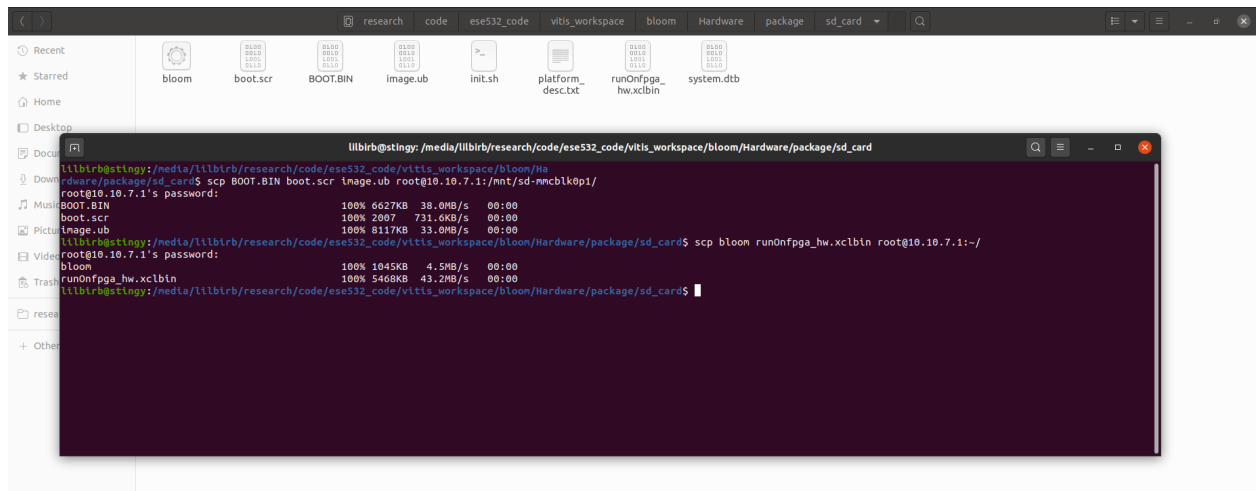
10. Click on `bloom.prj`. Check out the **Hardware optimization** option where you can change the optimization level for the hardware function. Additionally, recall from P2 that you can change the optimization level of the host code from the C/C++ build settings. Now click on the build button on the menu bar to start compilation:



11. Once the compilation completes, open the **Hardware** folder from the Explorer. The binaries are in the **package/sd_card** folder.



12. Copy the binaries and the **xrt.ini** to the Ultra96 as follows and then reboot the Ultra96.



13. Run the code using the following commands in the Ultra96:

```
ifconfig eth0 10.10.7.1 netmask 255.0.0.0
export XILINX_XRT=/usr
./bloom 40000 64
```

You should see the following output in the terminal:

```
root@ultra96v2-2020-1:~# ./bloom 40000 64
Initializing data
Creating documents - total size : 559.858 MBytes (139964416 words)
Creating profile weights
[ 1018.547572] [drm] Pid 769 opened device
[ 1018.551450] [drm] Pid 769 closed device
[ 1018.558627] [drm] Pid 769 opened device
Loading runOnfpga_hw.xclbin
[ 1018.617733] [drm] zocl_xclbin_read_axlf The XCLBIN already loaded
[ 1018.617765] [drm] zocl_xclbin_read_axlf 3c650f2f-9cc2-408a-8c92-0ec3bc33
[ 1018.633496] [drm] bitstream 3c650f2f-9cc2-408a-8c92-0ec3bc335ce3 locked,
[ 1018.641197] [drm] Reconfiguration not supported
[ 1018.652995] [drm] bitstream 3c650f2f-9cc2-408a-8c92-0ec3bc335ce3 unlocke
Processing 559.858 MBytes of data
Splitting data in 64 sub-buffers of 8.748 MBytes for FPGA processing
-----
[ooqueue]: Completed buffer migrate
[ooqueue]: Completed buffer migrate
[ooqueue]: Completed buffer migrate
[ooqueue]: Completed buffer migrate
[ooqueue]: Completed buffer migrate
[ooqueue]: Completed buffer migrate
```

[illegible]


•

•

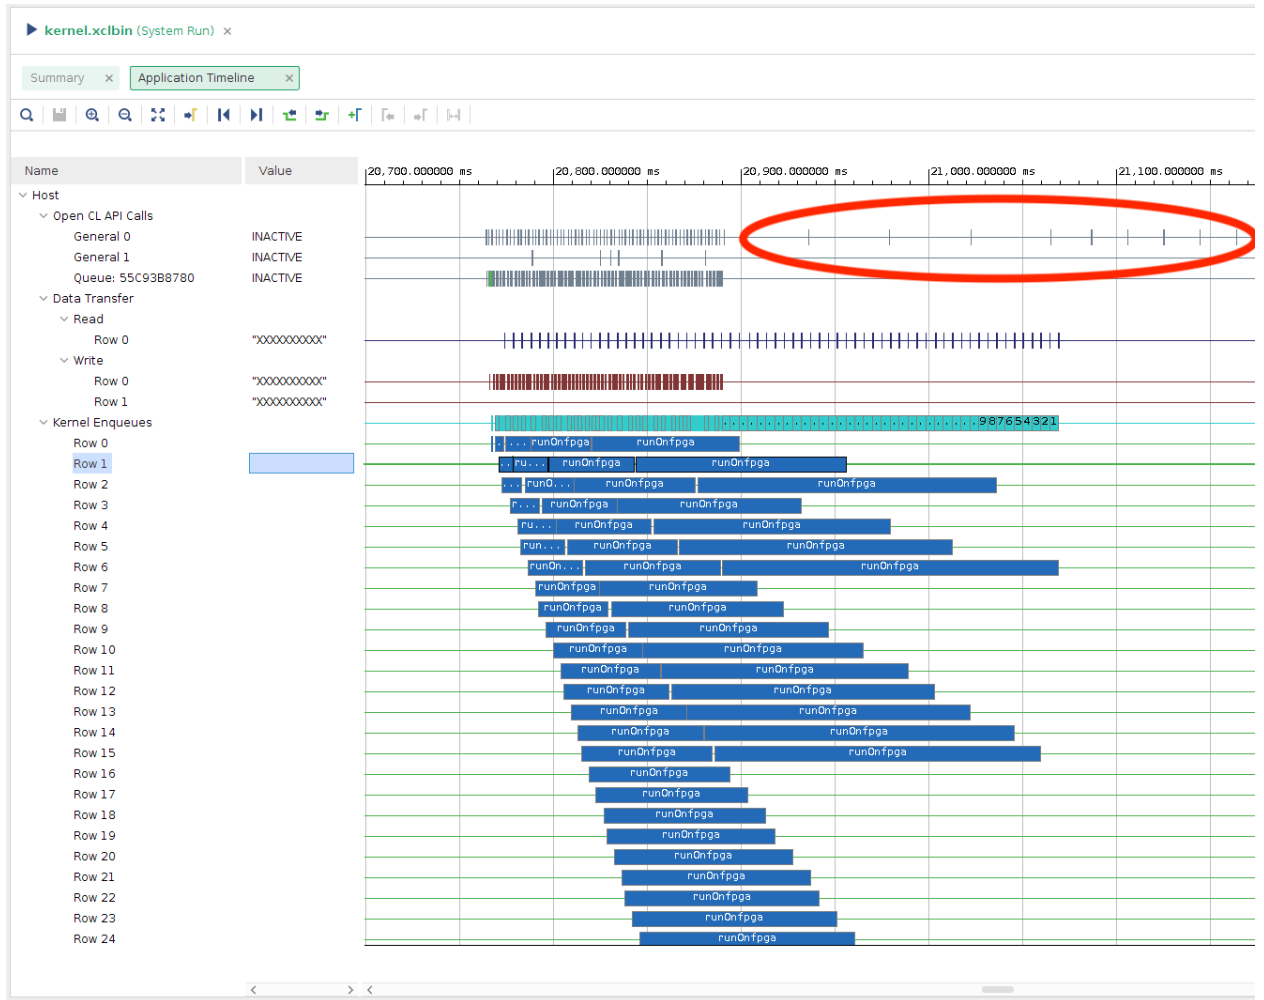
Executed FPGA accelerated version		1414.5707 ms	(FPGA 247.878 ms)
Executed Software-Only version		11779.4325 ms	

Verification: PASS

14. The verbose outputs in the terminal is caused by the call `set_callback(flagDone, "oooqueue");` in the host code. This is a helper function in `xc12.hpp` that prints out the state of an OpenCL event. You can use it to debug OpenCL calls. In addition, you can also use the `OCL_CHECK` macro from `xc12.hpp` to see if an OpenCL call succeeded.
15. Copy the generated run summary and csv files to your host computer and open vitis analyzer. You can see overlap of kernel execution with data transfer and OpenCL API calls.



16. If you scroll forward in the timeline, you can see overlap between computation in the cpu and the fpga as shown below.



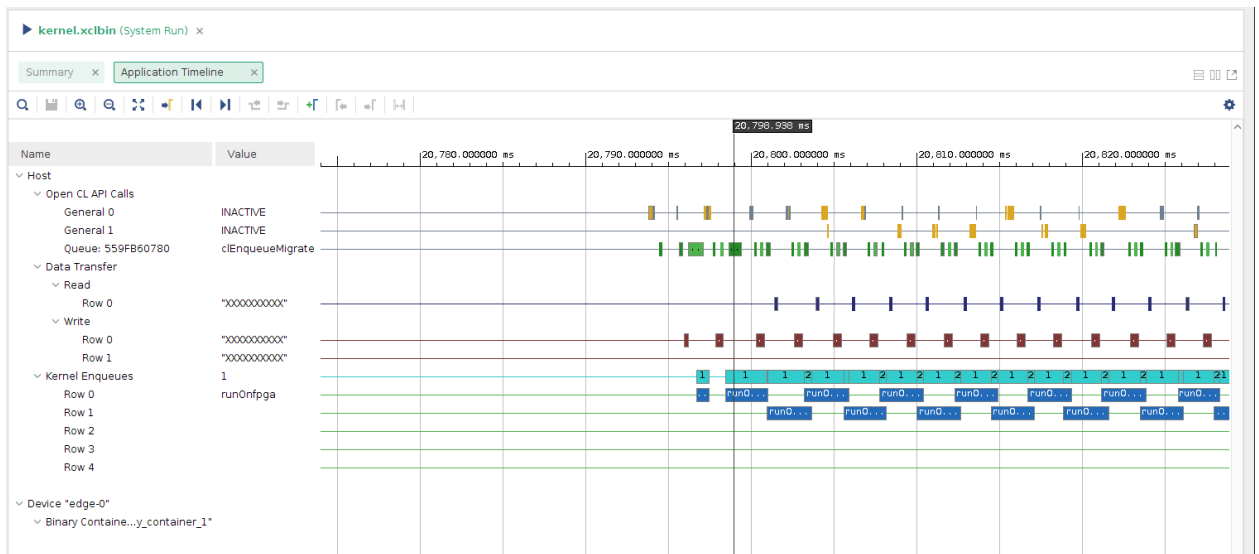
From the output in 13, these calls correspond to the `waiting...` print outs. You can check in the host code, how we wait for a `cl::Event` to finish based on a condition, and when the event notifies that it's finished, we start executing the cpu code, so that it overlaps with the fpga execution:

```
needed += size;
if (needed > available) {
    clWaitForEvents(1, (const cl_event *) &flagWait[iter]);

    std::cout << "waiting..." << std::endl;
    available += subbuf_doc_info[iter].size / sizeof(uint);
    iter++;
}
```

17. Now run with a different ITER value and look at the updated trace:

```
./bloom 40000 128
```

You can see that since the kernel execution time gets smaller as you increase the iteration number, the next kernel execution starts almost immediately.

18. Running a sweep on the number of iterations, we see that ITER=32 is the most performant for this design:

```
./bloom 40000 8
Executed FPGA accelerated version | 1413.3252 ms ( FPGA 305.796 ms )
Executed Software-Only version   | 11780.3703 ms
```

Verification: PASS

```
./bloom 40000 16
Executed FPGA accelerated version | 1396.5072 ms ( FPGA 298.034 ms )
Executed Software-Only version   | 11770.1858 ms
```

Verification: PASS

```
./bloom 40000 32
Executed FPGA accelerated version | 1391.2062 ms ( FPGA 284.336 ms )
Executed Software-Only version   | 11768.1306 ms
```

Verification: PASS

```
./bloom 40000 64
Executed FPGA accelerated version | 1414.5707 ms ( FPGA 247.878 ms )
Executed Software-Only version   | 11779.4325 ms
```

Verification: PASS

```
./bloom 40000 128
Executed FPGA accelerated version | 1458.1155 ms ( FPGA 179.195 ms )
Executed Software-Only version   | 11782.2403 ms
```

Verification: PASS

```
./bloom 40000 256
Executed FPGA accelerated version | 1533.5603 ms ( FPGA 10.701 ms )
Executed Software-Only version   | 11798.4502 ms
```

Verification: PASS

19. This concludes a top-down walk-through of this tutorial. To learn more about this design, read the following in-order:

- (a) [Overview of the Original Application](#)
- (b) [Architect a Device-Accelerated Application](#)
- (c) [Implementing the Kernel](#)
- (d) [Data Movement Between the Host and Kernel](#)

Note that the tutorial is written for data center cards. Some of the parameter choices, such as port data width, DDR memory etc. should be reconsidered for the Ultra96 to get optimal performance (refer to this paper: [Unexpected Diversity: Quantitative Memory Analysis for Zynq UltraScale+ Systems](#)).

Questions

If anything is unclear please post on Ed Discuss or come to office hours, and we will be glad to assist.